EVA Tech Stack

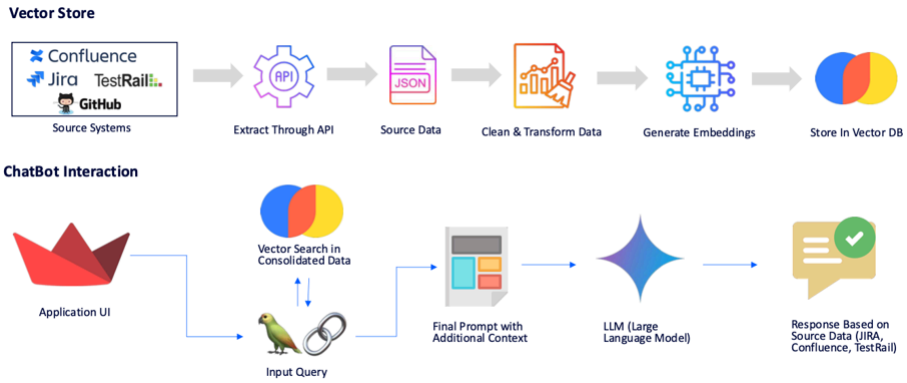

EVA, our Enterprise Chatbot, utilizes Retrieval Augmented Generation (RAG) for its operation. It extracts data from Confluence, Jira, TestRail, and Github using their respective API services. This data is then transformed and cleaned including custom fields, to ensure it's in the optimal format. The cleaned data is converted into a format that our AI can understand, a process known as 'Generate Embeddings', and is stored in the 'Vector DB'. When a user inputs a query, EVA performs a 'Vector Search in Consolidated Data' stored in the Vector DB.

Retrieval Augmented Generation (RAG)

The user query is converted into embeddings that the AI can understand, and a match is identified based on the meaning of the query using AI/ML algorithms. The prompt is filled with additional context and chat history, ensuring the query has relevant additional information from our internal tools. The 'Large Language Model' (LLM) processes this information to generate a response. The final output is a 'Response Based on Source Data' from Jira, Confluence, and TestRail. If these sources do not have relevant information, the LLM will attempt to answer with the general knowledge it has.